

CAR-PRICE-PREDICTION-PROJECT

SUBMITTED BY: - TEJENDRA SONI

ACKNOWLEDGMENT

I would like to express my thanks of gratitude to SME SWATI MAHASETH as well as Flip Robo who gave me the golden opportunity to do this wonderful project on the topic Micro Credit Defaulter, which also helped me in doing a lot of research and I came to know about so many new things. I am very thankful to them. I am making this project to increase my knowledge.

INTRODUCTION

Problem Statement:

With the covid 19 impact in the market, we have seen lot of changes in the car market. Now some cars are in demand hence making them costly and some are not in demand hence cheaper. One of our clients works with small traders, who sell used cars. With the change in market due to covid 19 impact, our client is facing problems with their previous car price valuation machine learning models. So, they are looking for new machine learning models from new data. We must make car price valuation mode

Objective:

Building a model which can be used to predict in terms of a probability for used cars

Firstly, we will start by importing required libraries and databases.

```
In [1]: #Importing needed libraries

import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import sklearn
import warnings
warnings.filterwarnings('ignore')
```

```
In [2]: #Importing the dataframe storing it into df & printing it
```

```
df=pd.read_csv('CAR-PRICE-DATA-FINAL.csv')
df
```

```
Out[2]:
```

	Unnamed: 0	Unnamed: 0.1	PRICE	YEAR	KM	COMPANY	LOCATION	POSTED-DATE	sourcefilename
0	0	0	645000.0	2018	41,250	Maruti Suzuki Baleno	KowdianJan 03	Jan-03	E:/TEJ/FLIPROBO/CAR-PRICE/DATAA/CARS-FILE01.csv
1	1	1	725000.0	2017	82,000	Maruti Suzuki Swift Dzire	Changubettyin5 days ago	5 days ago	E:/TEJ/FLIPROBO/CAR-PRICE/DATAA/CARS-FILE01.csv
2	2	2	13500000.0	2020	5,005	BMW 8 Series	CompanypadiJan 12	Jan-12	E:/TEJ/FLIPROBO/CAR-PRICE/DATAA/CARS-FILE01.csv
3	3	3	35000.0	1998	50,000	Maruti Suzuki 800	PazhakuttyToday	Today	E:/TEJ/FLIPROBO/CAR-PRICE/DATAA/CARS-FILE01.csv
4	4	4	185000.0	2008	1,41,000	Maruti Suzuki Swift	PalarivattomToday	Today	E:/TEJ/FLIPROBO/CAR-PRICE/DATAA/CARS-FILE01.csv
...
5492	495	495	1390000.0	2014	58,000	Audi Q3	Sector 80Today	Today	E:/TEJ/FLIPROBO/CAR-PRICE/DATAA/CARS-FILE9.csv
5493	496	496	280000.0	2012	80,000	Maruti Suzuki Ertiga	Gohana PartToday	Today	E:/TEJ/FLIPROBO/CAR-PRICE/DATAA/CARS-FILE9.csv
5494	497	497	45000.0	2006	1,18,000	Maruti Suzuki Alto 800	Prakash NagarToday	Today	E:/TEJ/FLIPROBO/CAR-PRICE/DATAA/CARS-FILE9.csv
5495	498	498	325000.0	2017	88,000	Hyundai Xcent	New Grain MarketToday	Today	E:/TEJ/FLIPROBO/CAR-PRICE/DATAA/CARS-FILE9.csv
5496	499	499	700000.0	2018	73,000	Hyundai i20 Active	Tagra Kali RamToday	Today	E:/TEJ/FLIPROBO/CAR-PRICE/DATAA/CARS-FILE9.csv

5497 rows x 9 columns

Our dataset has 5497 rows and 9 columns.

Let's check the datatype of all columns:

```
In [7]: #Printing the datatypes of all the columns that are presented here in our dataset
```

```
df.dtypes
```

```
Out[7]: Unnamed: 0      int64  
        Unnamed: 0.1    int64  
        PRICE          float64  
        YEAR           int64  
        KM             object  
        COMPANY        object  
        LOCATION       object  
        POSTED-DATE    object  
        soucefilename   object  
        dtype: object
```


- Also, we can drop column Unnamed: 0 & Unnamed: 0.1 as it contains serial number only.
- After dropping the columns, we don't need we can also see that there are only 2 columns are of integer type i.e., PRICE & YEAR

VISUALIZATIONS:

```
In [12]: #The white spots shows that there are null values in these columns
```

```
sns.heatmap(df.isnull())
```

```
Out[12]: <AxesSubplot:>
```




There are some null values present in our dataset showing it with the help of seaborn library & with the heatmap

All the white spots that you can see in the plot are null values shown here

```
In [41]: #Plotting catplot
```

```
sns.catplot(x='PRICE',y='YEAR',data=df)
```

```
Out[41]: <seaborn.axisgrid.FacetGrid at 0x2081b3583a0>
```




The categorical plot stats that PRICE & YEAR are in positive relationship we can here see as the year increases the Price of the cars are also increases

As a used car with low runner KM & latest as per MANUFACT-YEAR the PRICE is also increased for that

```
In [44]: #Plotting the relational plot of the PRICE & YEAR column with respect to the companies  
#The below shows us so many insights  
#Different company model shows their unique impact we can see it in the below graph  
  
sns.relplot(x='PRICE',y='YEAR',hue='COMPANY',data=df,kind='line')
```

```
Out[44]: <seaborn.axisgrid.FacetGrid at 0x2081b7fa2e0>
```

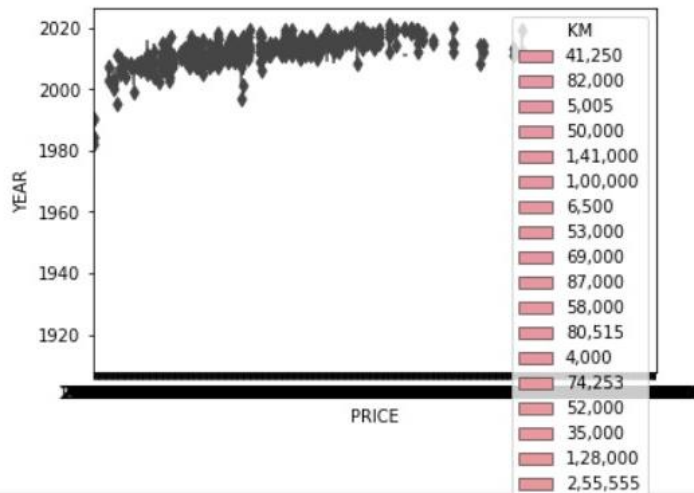


Here we have set hue=company means the plot will work with respect to company


```
In [45]: #Plotting boxen plot of PRICE & YEAR with respect to KM
#The KMS that are used per car are showed here with different colours

sns.boxenplot(x='PRICE',y='YEAR',hue='KM',data=df)
```

```
Out[45]: <AxesSubplot:xlabel='PRICE', ylabel='YEAR'>
```




With respect to KMS also we are here plotting the boxenplot

Different lines are representing the KM runner by the cars

In [47]: *#Plotting Jointplot*

```
sns.jointplot(x='PRICE',y='YEAR',data=df)
```

Out[47]: <seaborn.axisgrid.JointGrid at 0x20827e64a90>



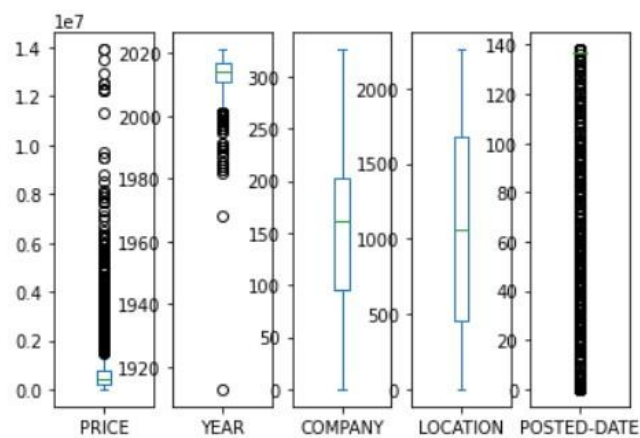
Plotting the joint plot of the variables

OUTLIERS:

```
In [50]: #Showing boxplots of all the variables
```

```
df.plot(kind='box',subplots=True)
```

```
Out[50]: PRICE      AxesSubplot(0.125,0.125;0.133621x0.755)  
YEAR      AxesSubplot(0.285345,0.125;0.133621x0.755)  
COMPANY   AxesSubplot(0.44569,0.125;0.133621x0.755)  
LOCATION   AxesSubplot(0.606034,0.125;0.133621x0.755)  
POSTED-DATE AxesSubplot(0.766379,0.125;0.133621x0.755)  
dtype: object
```




We can see outliers in many columns, we will check them & will remove them

LET'S CHECK THE SKEWNESS WITH THE HELP OF DISTRIBUTION PLOT

```
In [49]: #Determining the skewness present in all the columns  
#There is skewness present in our target variable
```

```
df.plot(kind='kde',subplots=True)
```

```
Out[49]: array([<AxesSubplot:ylabel='Density'>, <AxesSubplot:ylabel='Density'>,  
                <AxesSubplot:ylabel='Density'>, <AxesSubplot:ylabel='Density'>,  
                <AxesSubplot:ylabel='Density'>], dtype=object)
```



As we can clearly see there is some skewness present in PRICE

Also, in other columns as well

CORRELATION: -

```
In [53]: #Showing corelation through heatmap & also printing values for better understanding  
#Here all the DARK COLUMNS are HIGHLY CORELATED
```

```
sns.heatmap(dfcor,annot=True,cmap='Blues')
```

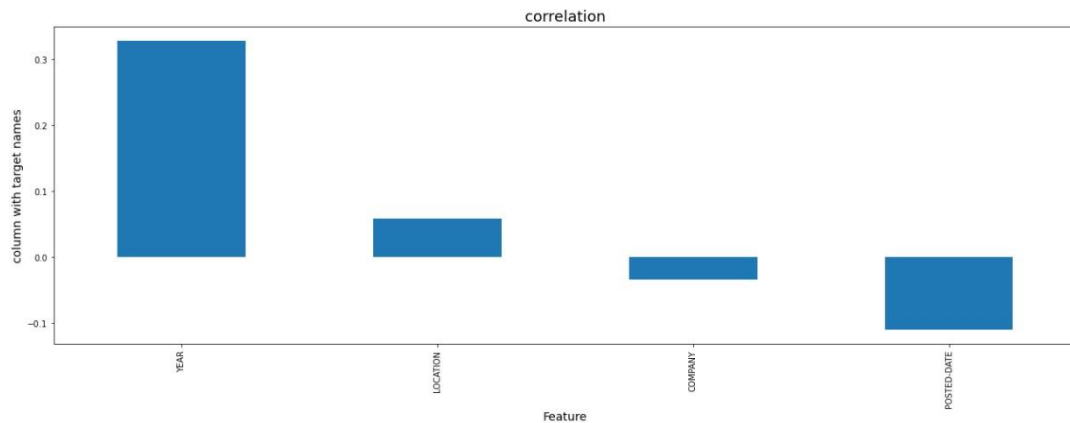
```
Out[53]: <AxesSubplot:>
```



All the dark blue shades show in the heatmap
Are highly co-related to each other

In [54]: *#Showing all the columns positively & negatively correlated*

```
plt.figure(figsize=(22,7))
df.corr()['PRICE'].sort_values(ascending=False).drop(['PRICE']).plot(kind='bar')
plt.xlabel('Feature',fontsize=14)
plt.ylabel('column with target names',fontsize=14)
plt.title('correlation',fontsize=18)
plt.show()
```




Only POSTED-DATE column is negatively correlated

We can see clearly in the above graph

PAIRPLOT: -

```
In [48]: #Plotting all the variables in respect to all the variable to determine each relationships b/w the variables
sns.pairplot(df)
```

```
Out[48]: <seaborn.axisgrid.PairGrid at 0x20829196bb0>
```



Plotting all the columns with respect to all the rows

OUTLIERS TREATMENT

```
In [61]: #Removing outliers

from scipy.stats import zscore
z=np.abs(zscore(df))
threshold=3
np.where(z>3)

Out[61]: (array([ 2, 3, 42, 148, 191, 220, 221, 266, 295, 303, 354,
365, 534, 553, 557, 580, 597, 630, 670, 693, 726, 766,
835, 849, 976, 977, 978, 979, 985, 987, 988, 989, 991,
1034, 1053, 1057, 1080, 1097, 1130, 1170, 1193, 1226, 1266, 1335,
1349, 1476, 1477, 1478, 1479, 1485, 1487, 1488, 1489, 1491, 1595,
1641, 1708, 1861, 1928, 2050, 2081, 2182, 2326, 2385, 2400, 2502,
2509, 2557, 2871, 2941, 2943, 2977, 3030, 3031, 3074, 3080, 3092,
3108, 3191, 3299, 3416, 3470, 3472, 3491, 3492, 3512, 3524, 3563,
3564, 3582, 3591, 3600, 3634, 3682, 3710, 3741, 3742, 3743, 3744,
3753, 3754, 3769, 3785, 3786, 3846, 3859, 3871, 3901, 3902, 3904,
3905, 3918, 3969, 3991, 4001, 4063, 4083, 4143, 4265, 4326, 4483,
4534, 4537, 4544, 4549, 4562, 4565, 4630, 4651, 4670, 4685, 4690,
4694, 4750, 4754, 4785, 4851, 4930, 4935, 4949, 4957], dtype=int64),
array([0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 1, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1, 0, 0, 1, 1,
0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0], dtype=int64))
```

```
In [62]: #Storing the dataframe after removing outlier
```

```
df_new=df[(z<3)].all(axis=1)
df_new
```

```
Out[62]:
```

	PRICE	YEAR	COMPANY	LOCATION	POSTED-DATE
0	645000.0	2018	183	1034	48
1	725000.0	2017	184	360	3
4	105000.0	2008	183	1497	137
5	125000.0	2006	160	1700	137
6	145000.0	2015	283	1502	137
...
5452	1390000.0	2014	8	1862	137
5493	280000.0	2012	170	622	137
5494	45000.0	2006	161	1605	137
5495	325000.0	2017	104	1413	137
5496	700000.0	2018	107	2010	137

4828 rows x 5 columns

Removing outliers using z-score technique

Saving the cleaned data into the new data frame

As we can see the shape has also changed to 4282 rows

& 5 columns

SKEWNESS REMOVING: -

```
In [67]: #Removnig skewness with the help of power transformation
```

```
from sklearn.preprocessing import power_transform  
df_new=power_transform(x)  
  
df_new=pd.DataFrame(df_new,columns=x.columns)
```

```
In [68]: #Here we can see skewness has been removed from the data
```

```
df_new.skew()
```

```
Out[68]: YEAR          -0.210969  
COMPANY          -0.103469  
LOCATION          -0.250172  
POSTED-DATE     -1.431037  
dtype: float64
```

We have removed the skewness to the maximum extent we can using power transformation technique

FINDING BEST RANDOM STATE:

In [70]: *#Using for loop determining the best accuracy for the model at the best random state*

```
maxAccu=0
maxRS=0
for i in range(1,100):
    x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=.10,random_state=i)
    DTC=DecisionTreeClassifier()
    DTC.fit(x_train,y_train)
    pred=DTC.predict(x_test)
    acc=accuracy_score(y_test,pred)
    if acc>maxAccu:
        maxAccu=acc
        maxRS=i
print("Best accuracy is ",maxAccu, " on Random State ",maxRS)
```

Best accuracy is 0.2836438923395445 on Random State 8

With the help of for loop we are determining the best possible accuracy at the best possible random state

MODEL/S DEVELOPMENT AND EVALUATION:

```
In [71]: #Sending the data for training & testing phase  
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=.10,random_state=maxRS)
```

```
In [72]: #Created a list in which we have stored all the instances of the model  
#Using for loop we will determine the accuracy of all the models  
  
model=[DecisionTreeClassifier(),SVC(),AdaBoostClassifier(),RandomForestClassifier(),LogisticRegression()]  
  
for m in model:  
    m.fit(x_train,y_train)  
    #m.score(x_train,y_train)  
    pred=m.predict(x_test)  
    acc=accuracy_score(y_test,pred)  
    print('Accuracy Score of',m,'is:', acc)  
    print(confusion_matrix(y_test,pred))  
    print(classification_report(y_test,pred))  
    print('\n')
```

Sending the data to training & testing phase with the help of train test split

Again, with the help of for loop we are here determining the accuracies of all the models for selecting the best model for our dataset

MODEL SELECTION

MODELS	ACCCURACY
DTC	28.57
SVC	00.62
ABC	00.14
RFC	27.32
LR	00.08

Here in the table, we can see that DTC model is giving

Highest accuracy among all the different present models

So, we select the DTC best for our model

But before proceeding further we will cross validate

Each of the models so that we can finally select which model will be the best for our dataset

CROSS VALIDATION: -

```
In [90]: #Again using for loop we will now determine the cross validation of each model

model=[DecisionTreeClassifier(),SVC(),AdaBoostClassifier(),RandomForestClassifier(),LogisticRegression()]
for m in model:
    score=cross_val_score(m,x,y,cv=5)
    print("Score for ",m," is : ",score.mean())

Score for DecisionTreeClassifier() is : 0.21226702710820752
Score for SVC() is : 0.012220255527306666
Score for AdaBoostClassifier() is : 0.017190916014975487
Score for RandomForestClassifier() is : 0.21267981849193834
Score for LogisticRegression() is : 0.007869640309378988
```

After determining the cross-validation scores of the models we will now select RFC as our model

Because the RFC model shows the least difference in all the models

HYPER PARAMETER TUNNING: -

```
In [106]: #Importing the grid search cv for hypertune the model
          from sklearn.model_selection import GridSearchCV

In [107]: #Passing the parameters for the model
          parameters = {'max_depth':np.arange(2,10),
                        'criterion':['gini','entropy'],
                        'max_features':['auto','sqrt','log2'],
                        'min_samples_split':[2,3,4]}

In [108]: #Creating the instance of the model
          GCV=GridSearchCV(RandomForestClassifier(),parameters,cv=5)

In [109]: #Fetting the model
          GCV.fit(x_train,y_train)

Out[109]: GridSearchCV(cv=5, estimator=RandomForestClassifier(),
                    param_grid={'criterion': ['gini', 'entropy'],
                                'max_depth': array([2, 3, 4, 5, 6, 7, 8, 9]),
                                'max_features': ['auto', 'sqrt', 'log2'],
                                'min_samples_split': [2, 3, 4]})

In [110]: #Getting the best parameters
          GCV.best_params_

Out[110]: {'criterion': 'entropy',
           'max_depth': 9,
           'max_features': 'sqrt',
           'min_samples_split': 3}

In [111]: #Passing the best parameters & printing the final accuracy
          Final_mod= RandomForestClassifier(criterion="entropy",max_depth=9,max_features="sqrt",min_samples_split=3)
          Final_mod.fit(x_train,y_train)
          pred=Final_mod.predict(x_test)
          acc=accuracy_score(y_test,pred)
          print(acc*100)


23.60248447204969
```

Tunning our model to get the best accuracy of the model

PLOTTING A GRAPH STATING THE PREDICTED & MEASURED RESULTS

In [113]: *#Ploting a diagram of the predicted & measured results*

```
fig_dims = (20, 10)
fig, ax = plt.subplots(figsize=fig_dims)
ax.scatter(y_test, pred)
ax.plot([y_test.min(), y_test.max()], [y_test.min(), y_test.max()], 'k--', lw=4)
ax.set_xlabel('Measured')
ax.set_ylabel('Predicted')
plt.show()
```



The line that tries to cover the Blues Points that is our predicted Line Which tries to get mix with the Measured results

SAVING THE MODEL: -

```
In [114]: #Importing pickle for saving the model  
#Saving it in the pickle file  
  
import pickle  
filename= 'CAR-PREDICTION.pkl'  
pickle.dump(Final_mod, open(filename, 'wb'))
```

Saving the model with the help of pickle library

CONCLUSION: -

```
In [115]: #Load the model from the disk

loaded_model = pickle.load(open('CAR-PREDICTION.pkl', 'rb'))
result = loaded_model.score(x_test,y_test)
print(result)

0.2360248447204969

In [116]: #Printing the dataframe of the predicted & measured

conclusion=pd.DataFrame([loaded_model.predict(x_test)[:],pred[:]],index=["Predicted","Original"])
conclusion

Out[116]:
```

	0	1	2	3	4	5	6	7	8	9	...	473	474	475	476
Predicted	725000.0	625000.0	265000.0	575000.0	200000.0	3600000.0	510000.0	750000.0	3100000.0	835000.0	...	150000.0	3999999.0	485000.0	640000.0
Original	725000.0	625000.0	265000.0	575000.0	200000.0	3600000.0	510000.0	750000.0	3100000.0	835000.0	...	150000.0	3999999.0	485000.0	640000.0

2 rows x 483 columns

```
In [117]: #END
```

Also, making a data frame of predicted answers & measured answers & printing it